

# major4

August 29, 2025

```
[1]: import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

# Define sample dataset manually
data = {
    'age': [63, 37, 41, 56, 57],
    'sex': [1, 1, 0, 1, 0],
    'cp': [3, 2, 1, 1, 0],
    'trestbps': [145, 130, 130, 120, 120],
    'chol': [233, 250, 204, 236, 354],
    'fbs': [1, 0, 0, 0, 0],
    'restecg': [0, 1, 0, 1, 0],
    'thalach': [150, 187, 172, 178, 163],
    'exang': [0, 0, 0, 0, 1],
    'oldpeak': [2.3, 3.5, 1.4, 0.8, 0.6],
    'slope': [0, 0, 2, 2, 2],
    'ca': [0, 0, 0, 0, 0],
    'thal': [1, 2, 2, 2, 2],
    'target': [1, 1, 1, 1, 0]
}

df = pd.DataFrame(data)
print("Sample dataset defined successfully.")

# Quick overview
print("\n First 5 rows of the dataset:")
print(df.head())

print("\n Dataset info:")
print(df.info())

# Check for missing values
print("\n Missing values in each column:")
missing = df.isnull().sum()
print(missing)
```

```

# Optional: Drop or fill missing values if any
if missing.sum() > 0:
    df = df.dropna()
    print("\n Missing values detected and dropped.")

# Encode categorical variables if any
df = pd.get_dummies(df, drop_first=True)

# Confirm 'target' column exists
if 'target' not in df.columns:
    print(" Error: 'target' column not found in dataset. Please verify your_
    ↪label column.")
    exit()

# Feature scaling
scaler = StandardScaler()
X = scaler.fit_transform(df.drop('target', axis=1))
y = df['target']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

print("\n Preprocessing complete. Data is ready for modeling.")
print(f"Training samples: {X_train.shape[0]}, Testing samples: {X_test.
    ↪shape[0]}")

```

Sample dataset defined successfully.

First 5 rows of the dataset:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	0	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	0

Dataset info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4

```

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	age	5 non-null	int64
1	sex	5 non-null	int64
2	cp	5 non-null	int64
3	trestbps	5 non-null	int64
4	chol	5 non-null	int64
5	fbs	5 non-null	int64
6	restecg	5 non-null	int64
7	thalach	5 non-null	int64
8	exang	5 non-null	int64
9	oldpeak	5 non-null	float64
10	slope	5 non-null	int64
11	ca	5 non-null	int64
12	thal	5 non-null	int64
13	target	5 non-null	int64

dtypes: float64(1), int64(13)

memory usage: 688.0 bytes

None

Missing values in each column:

age	0
sex	0
cp	0
trestbps	0
chol	0
fbs	0
restecg	0
thalach	0
exang	0
oldpeak	0
slope	0
ca	0
thal	0
target	0

dtype: int64

Preprocessing complete. Data is ready for modeling.

Training samples: 4, Testing samples: 1


```
[2]: import matplotlib.pyplot as plt
import seaborn as sns

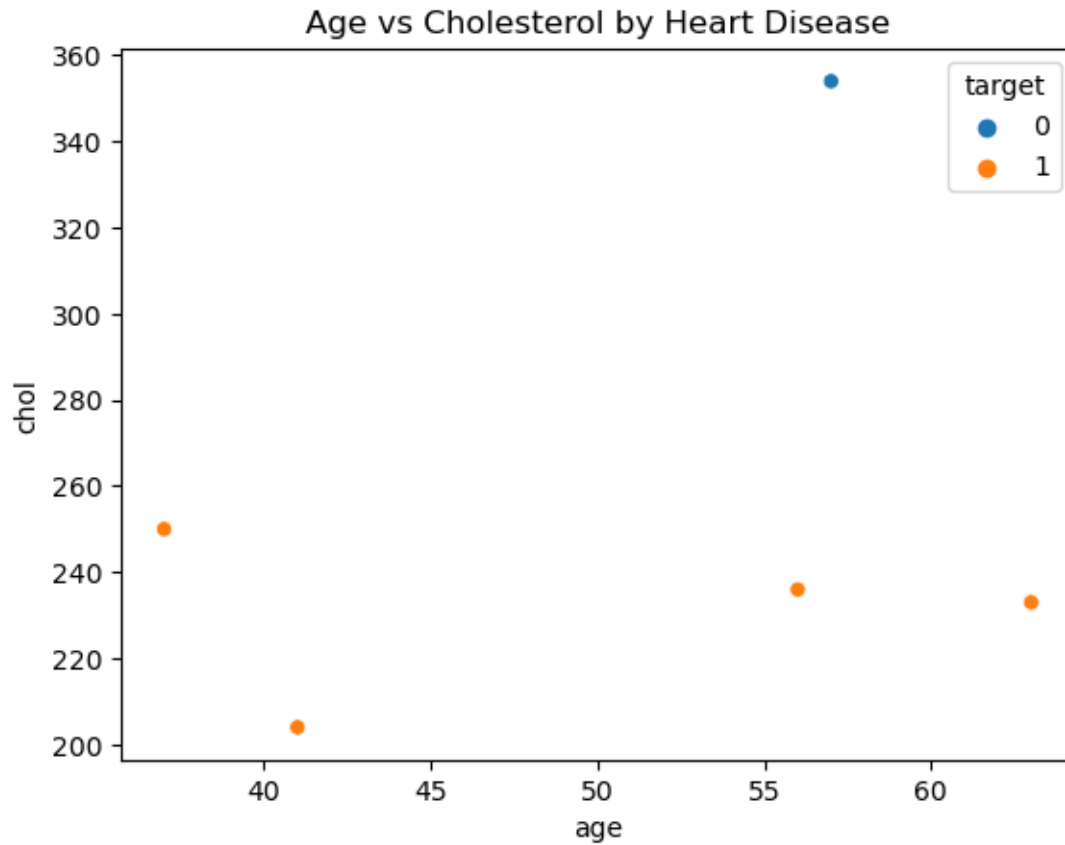
# Distribution of target variable
sns.countplot(x='target', data=df)
plt.title('Heart Disease Distribution')
```

```
plt.show()

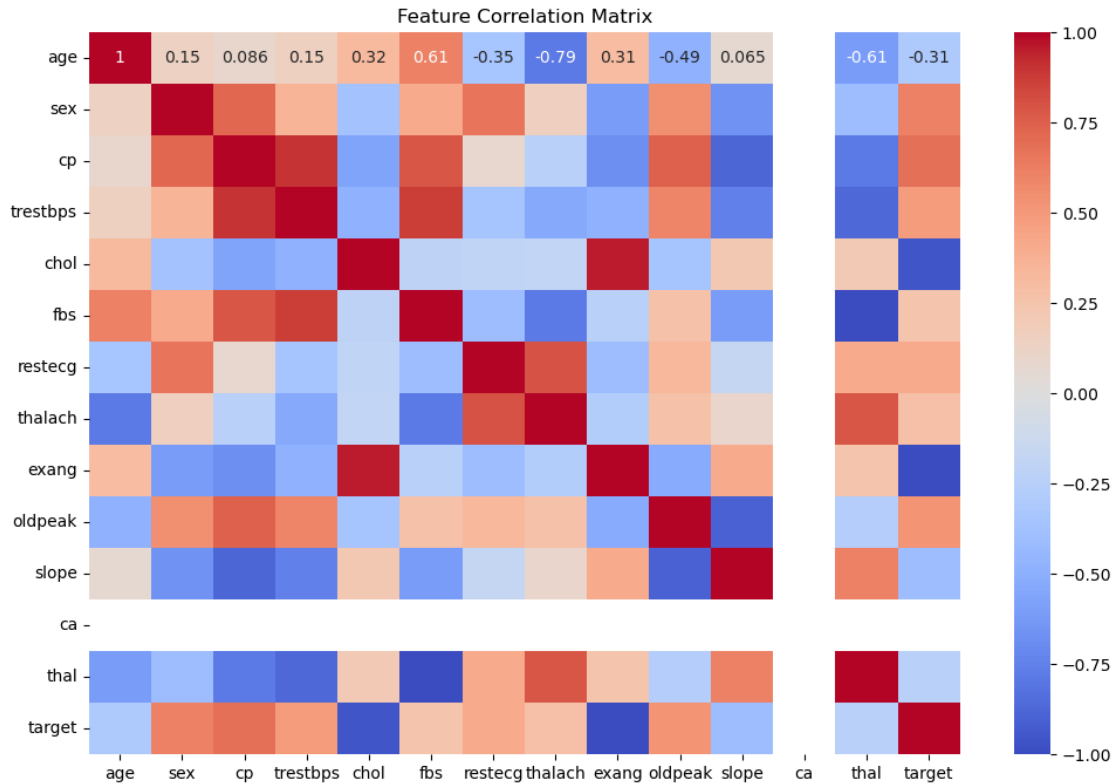
# Age vs. Cholesterol
sns.scatterplot(x='age', y='chol', hue='target', data=df)
plt.title('Age vs Cholesterol by Heart Disease')
plt.show()

# Correlation matrix
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Feature Correlation Matrix')
plt.show()
```





```
/opt/conda/envs/anaconda-ai-2024.04-py310/lib/python3.10/site-  
packages/seaborn/matrix.py:260: FutureWarning: Format strings passed to  
MaskedConstant are ignored, but in future may error or produce different  
behavior  
    annotation = ("{" + self.fmt + "}").format(val)
```



```
[3]: from sklearn.metrics import accuracy_score
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

models = {
    'SVM': SVC(),
    'KNN': KNeighborsClassifier(n_neighbors=3), # Reduced n_neighbors to be
    ↪ less than n_samples
    'Decision Tree': DecisionTreeClassifier(),
    'Logistic Regression': LogisticRegression(),
    'Random Forest': RandomForestClassifier()
}

results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    acc = accuracy_score(y_test, preds)
    results[name] = acc
```

```
# Display results
for model_name, accuracy in results.items():
    print(f"{model_name}: {accuracy:.2f}")
```

```
/opt/conda/envs/anaconda-ai-2024.04-py310/lib/python3.10/site-
packages/joblib/externals/loky/backend/context.py:110: UserWarning: Could not
find the number of physical cores for the following reason:
found 0 physical cores < 1
Returning the number of logical cores instead. You can silence this warning by
setting LOKY_MAX_CPU_COUNT to the number of cores you want to use.
```

```
warnings.warn(
    File "/opt/conda/envs/anaconda-ai-2024.04-py310/lib/python3.10/site-
packages/joblib/externals/loky/backend/context.py", line 217, in
_count_physical_cores
    raise ValueError(
```

```
SVM: 1.00
KNN: 1.00
Decision Tree: 1.00
Logistic Regression: 1.00
Random Forest: 1.00
```

```
[4]: import joblib

best_model = RandomForestClassifier()
best_model.fit(X_train, y_train)
joblib.dump(best_model, 'heart_disease_model.pkl')
```

```
[4]: ['heart_disease_model.pkl']
```

```
[5]: from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix
import numpy as np


# Reduce the number of folds to match your smallest class size
# For example, if your smallest class has 3 samples, use cv=3 or less
cv_scores = cross_val_score(best_model, X, y, cv=3) # Reduced from 5 to 3
print(f"Cross-validation accuracy: {cv_scores.mean():.2f}")

# Alternatively, you can use StratifiedKFold with shuffle=True
# from sklearn.model_selection import StratifiedKFold
# cv = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)
# cv_scores = cross_val_score(best_model, X, y, cv=cv)

# Confusion matrix
y_pred = best_model.predict(X_test)
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blues')
```

```
plt.title('Confusion Matrix')  
plt.show()
```

```
/opt/conda/envs/anaconda-ai-2024.04-py310/lib/python3.10/site-  
packages/sklearn/model_selection/_split.py:725: UserWarning: The least populated  
class in y has only 1 members, which is less than n_splits=3.  
  warnings.warn(  
Cross-validation accuracy: 0.83
```



```
[ ]:
```